

✓ DATA PREPARATION

- Considering the CIFAR - 10 dataset
- Converting this dataset into binary dataset [which have label as 0 or 1]
- Converting this dataset into unbalanced dataset by converting labels '0' as '0' and rest everything label as '1'

```
1 from torchvision.datasets import CIFAR10
2 import torchvision.transforms as transforms
3 import os
4 import pandas as pd
5 from torchvision.transforms import ToPILImage
6
7 transform = transforms.ToTensor()
8
9 train_set = CIFAR10(root='./data', train=True, download=True)
10 test_set = CIFAR10(root='./data', train=False, download=True)
11
```

⇌ 100%|██████████| 170M/170M [00:15<00:00, 11.2MB/s]

```
1 def label_adjustment(dataset):
2     results = []
3     for img, label in dataset:
4         if label == 0:
5             results.append([img, label])
6         if label == 1 or label == 2:
7             results.append([img, 1])
8     return results
```

```
1 train_data = label_adjustment(train_set)
2 test_data = label_adjustment(test_set)
```

```
1 save_root = '/content/drive/MyDrive/cifar10_binary'
2 os.makedirs(save_root, exist_ok = True)
3 to_pil = ToPILImage()
4
5 def save_images_and_make_csv(data, split_name):
6     dir = os.path.join(save_root, split_name)
7     os.makedirs(dir, exist_ok = True)
8     rows = []
```

```

9     for idx, (img_tensor, label) in enumerate(data):
10         img_path = os.path.join(dir, f'{idx}.png')
11         to_pil(img_tensor).save(img_path)
12         rows.append([img_path, label])
13
14     df = pd.DataFrame(rows, columns=["MD5HASH", "LABEL"])
15     df.to_csv(os.path.join(save_root, f"{split_name}.csv"),
16             print(f"{split_name}.csv saved with {len(rows)} entries
17
18 # Save both splits
19 save_images_and_make_csv(train_data, "train")
20 save_images_and_make_csv(test_data, "test")
21

```

⇒ train.csv saved with 15000 entries.
test.csv saved with 3000 entries.

```
1 !ls /content/drive/MyDrive/cifar10_binary
```

⇒ test test.csv train train.csv

IMPORTING LIBRARIES AND SETTING CONGIURATION PARAMETERS

```

1 from torch.utils.data import Dataset, DataLoader
2 import torch
3 import torch.nn as nn
4 import torch.nn.functional as F
5 from sklearn.metrics import classification_report
6 from transformers import AutoImageProcessor, AutoModelForIm
7 from torch.optim.lr_scheduler import StepLR
8 # note: PIL stands for pillow; to install type "pip3 instal
9 from PIL import Image
10 from datetime import datetime
11 from torchvision.transforms import Compose, Resize, RandomR
12 import pandas as pd
13 import numpy as np
14 import os, gc

```

⇒ 2025-06-19 07:17:03.011678: E external/local_xla/xla/stream_executor/cuda/cud
WARNING: All log messages before absl::InitializeLog() is called are written
E0000 00:00:1750317423.209937 70 cuda_dnn.cc:8310] Unable to register cu
E0000 00:00:1750317423.267046 70 cuda_blas.cc:1418] Unable to register c

```

1 # Check for CUDA and MPS availability, set the device accor
2 if torch.backends.mps.is_available():
3     device = torch.device("mps")
4     # setting environment variables, need to run training i
5     os.environ['PYTORCH_MPS_HIGH_WATERMARK_RATIO'] = '0.0'
6     os.environ['PYDEVD_DISABLE_FILE_VALIDATION'] = '1'
7     print("Using MPS as the device.")
8 else:
9     if torch.cuda.is_available():
10    # the syntax 'cuda:3' used to point a specific GPU from
11    # 'cuda' points to first GPU from the cluster of GPUs
12        device = torch.device("cuda")
13        print("Using CUDA as the device.")
14    else:
15        device = torch.device("cpu")
16        print("Using CPU as the device.")

```

➞ Using CUDA as the device.

```

1 TRAINING_DATA = '/content/drive/MyD
2 TESTING_DATA = '/content/drive/MyD
3 BATCH_SIZE = 32#256
4 WORKERS = 4
5 PIN_MEMORY = True
6 MIXING = True
7 MODEL_NAME = "facebook/dinov2-ba
8 RESULTS = 'results'
9 EPOCHS = 4
10 BEST_MODEL = None
11 PRETRAINING = False
12 LEARNING_RATE = 5e-5
13 L2_PENALTY = 1e-5
14 GAMMA = 0.1
15 STEPSIZE = 3
16 SAVE_CHECKPOINTS = True
17 MIN_LOSS = float('inf')
18 MODEL_SAVED = f'{RESULTS}/bestmod
19 THRESHOLD = 0.5
20 OUTPUT_DIM = 1
21

```

✓ PERFORMANCE METRIC [PRECISION & RECALL]

```

1 def calculate_classification_accuracy(loader, model):
2     model.eval() # Set the model in evaluation mode
3     LABELS = []
4     PREDICTIONS = []
5
6     with torch.no_grad():
7         for images, labels in loader:
8             # Move to device and cast to float32
9             images, labels = images.to(device), labels.to(device)
10            probabilities = model(images).squeeze()
11            # Predictions based on the threshold
12            prediction = torch.where(probabilities > THRESHOLD, 1, 0)
13            LABELS.extend(labels.tolist())
14            PREDICTIONS.extend(prediction.tolist())
15    return classification_report(LABELS, PREDICTIONS)
16
17

```

DINO MODEL [WITH ADDITIONAL CLASSIFICATION HEAD]

```

1 class Network(nn.Module):
2     def __init__(self):
3         super(Network, self).__init__()
4
5         # taking processor for necessary substitutions, if needed
6         self.processor = AutoImageProcessor.from_pretrained('facebook/dinov2')
7         self.model = Dinov2Model.from_pretrained(MODEL_NAME)
8         self.pretrained_model_last_dim = self.model.layer_norm.weight.shape[0]
9
10        # Additional classification head used for downstream tasks
11        self.cls_head = nn.Sequential(
12            nn.Linear(self.pretrained_model_last_dim, OUTPUT_DIM),
13            nn.Sigmoid(),
14        )
15
16        self.initialize_weights() # weights initialization
17
18    def initialize_weights(self):
19        torch.manual_seed(444)
20        for layer in self.cls_head:
21            if isinstance(layer, nn.Linear):
22                nn.init.kaiming_uniform_(layer.weight, non_blocking=True)
23                nn.init.zeros_(layer.bias)

```

```
24         print(f"kaiming_uniform_ Initialization: {`  
25  
26  
27     def forward(self, x):  
28         x = self.model(x).last_hidden_state[:, 0]  
29         x = self.cls_head(x)  
30         return x  
31  
32
```

```
1 model = Network()  
2  
3 model.to(device)  
4  
5
```



```

preprocessor_config.json: 0%|          | 0.00/436 [00:00<?, ?B/s]
Using a slow image processor as `use_fast` is unset and a slow processor was
config.json: 0%|          | 0.00/548 [00:00<?, ?B/s]
model.safetensors: 0%|          | 0.00/346M [00:00<?, ?B/s]
kaiming_uniform_ Initialization: Linear
Network(
  (model): Dinov2Model(
    (embeddings): Dinov2Embeddings(
      (patch_embeddings): Dinov2PatchEmbeddings(
        (projection): Conv2d(3, 768, kernel_size=(14, 14), stride=(14, 14))
      )
      (dropout): Dropout(p=0.0, inplace=False)
    )
    (encoder): Dinov2Encoder(
      (layer): ModuleList(
        (0) Conv2d(768, 768, kernel_size=(14, 14), stride=(14, 14))
        (1) LayerNorm((768,))
        (2) Dinov2Attention(
          (norm1): LayerNorm((768,)), eps=1e-06, elementwise_affine=True)
          (attention): Dinov2Attention(

```

DATA LOADER WITH UPSAMPLING

```

1
2 class MD5HASHDataset(Dataset):
3     def __init__(self, dataframe):
4         self.dataframe = dataframe
5         self.images = self.dataframe['MD5HASH'].values
6         self.labels = self.dataframe['LABEL'].values
7         self.processor = model.processor
8         self.mean = self.processor.image_mean
9         self.std = self.processor.image_std
10        self.interpolation = self.processor.resample
11
12        self.train_transform = Compose([
13            Resize(size = (32, 32)),
14            #Resize(size = (85, 550)),
15            #RandomResizedCrop(size = (224, 224),
16            #                    scale = (0.08, 1.0),
17            #                    ratio = (0.75, 1.3333),
18            #                    interpolation = self.interpo
19            #RandomHorizontalFlip(p = 0.5),
20            #ColorJitter(brightness = (0.6, 1.4),
21            #            contrast = (0.6, 1.4),
22            #            saturation = (0.6, 1.4)),
23            ToTensor(),
24            Normalize(mean = self.mean, std = self.std),
25        ])
26
27
28    def __len__(self):
29        return len(self.images)
30
31    def __getitem__(self, idx):

```

```
32         # Load the image from the file path
33         image_path = self.images[idx]
34         image = self.train_transform(Image.open(image_path))
35         # Get the label
36         label = torch.tensor(self.labels[idx], dtype=torch.
37                               float)
38         return image, label
39
40
41 def create_training_loader(data_csv = TRAINING_DATA, upsamp
42                             # Load data
43                             training_data = pd.read_csv(data_csv)
44                             print('::: TRAINING DATA DETAILS :::')
45                             print('- Number of Samples:', training_data.shape[0])
46                             print('- LABEL DISTRIBUTION: \n', training_data['LABEL'])
47
48                             # Create dataset and dataloader
49                             md5hash_dataset = MD5HASHDataset(training_data)
50                             if upsampling:
51                                 # References:
52                                 # https://pytorch.org/docs/stable/data.html
53                                 # https://towardsdatascience.com/demystifying-pytor
54                                 from torch.utils.data import WeightedRandomSampler
55                                 classes_count = dict(training_data['LABEL'].value_c
56                                                      ounts.items())
57                                 sample_weights = [ 1 / classes_count[i] for i in tr
58                                                     aining_data['LABEL'].value_c
59                                                      ounts.keys()]
60                                 sampler = WeightedRandomSampler(weights = sample_we
61                                                                 ghts,
62                                                                 num_samples = len(t
63                                                                 raining_data['LABEL'].value_c
64                                                                 ounts.keys()),
65                                                                 replacement = True)
66                                 data_loader = DataLoader(md5hash_dataset,
67                                                         batch_size = BATCH_SIZE,
68                                                         num_workers = WORKERS,
69                                                         pin_memory = PIN_MEMORY,
70                                                         shuffle = False,
71                                                         sampler = sampler)
72                             else:
73                                 data_loader = DataLoader(md5hash_dataset,
74                                                         batch_size = BATCH_SIZE,
75                                                         num_workers = WORKERS,
76                                                         pin_memory = PIN_MEMORY,
77                                                         shuffle = MIXING)
78
79                             # Clean memory, :)
80                             del training_data
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
```

```

36     return data_loader
37
38

```

```

1 data_loader = create_training_loader(upsampling = True)
2

```

```

→ ::: TRAINING DATA DETAILS :::
- Number of Samples: 15000
- LABEL DISTRIBUTION:
  LABEL
1      10000
0       5000
Name: count, dtype: int64

```

✓ OPTIMIZER, SCHEDULER AND LOSS FUNCTION

```

1 # ## MODEL TRAINING
2
3 # Define the optimizer
4 # Idea borrowed from Research paper titled as "Improving Ge
5 if PRETRAINING:
6     optimizer = torch.optim.SGD(model.parameters(), lr = LE
7 else:
8     optimizer = torch.optim.Adam(model.parameters(), lr = L
9
10 # Define a learning rate scheduler
11 scheduler = StepLR(optimizer, step_size = STEPSIZE, gamma =
12 # Define the loss function: BCE
13 criterion = nn.BCELoss()
14
15

```

```

1 # directory creation
2 os.makedirs(RESULTS, exist_ok = True)
3 if SAVE_CHECKPOINTS:
4     CHECKPOINTDIR = f'{RESULTS}/checkpoints'
5     os.makedirs(CHECKPOINTDIR, exist_ok = True)
6

```

✓ MODEL FINE-TUNING TRAINING


```
1 # TRAINING LOOP
2 for epoch in range(EPOCHS):
3     print('-'*70)
4     # Define the total number of batches in the loader
5     total_loss = 0.0
6
7     # setting model stage to training
8     model.train()
9
10    for batch_idx, (images, labels) in enumerate(data_loader):
11        # shifting to MPS
12        # Shift to MPS and then cast to float32
13        images, labels = images.to(device), labels.to(device)
14
15        # Forward pass
16        optimizer.zero_grad() # Moved this line here to avoid
17
18        with torch.set_grad_enabled(True):
19            # Forward pass
20            outputs = model(images).squeeze() # Squeeze to 1D
21            loss = criterion(outputs, labels)
22
23            loss.backward()
24            optimizer.step()
25
26            total_loss += loss.item()
27            # Explicitly free up GPU memory
28            if torch.backends.mps.is_available():
29                torch.backends.mps.is_macos13_or_newer.cache_clear()
30            if torch.cuda.is_available():
31                torch.cuda.empty_cache()
32            # Run garbage collector to free up CPU memory
33            gc.collect()
34
35    print(f"Epoch {epoch + 1}/{EPOCHS}, Loss: {total_loss / len(data_loader)}")
36    print('TRAINING DATA')
37    print(f'- Performance: \n{calculate_classification_accuracy(model, data_loader)}')
38    # Update the learning rate
39    scheduler.step()
40
41
42    if SAVE_CHECKPOINTS:
43        timestamp = datetime.now().strftime("%Y%m%d%H%M%S")
44        checkpointmodel = '{}/epoch_{}_{}.pth'.format(CHECKPOINT_PATH, epoch + 1, timestamp)
45        print('Saving checkpoint: ', checkpointmodel)
```

```

46         torch.save(model.state_dict(), checkpointmodel)
47
48     # Check if this epoch had the minimum loss
49     if total_loss < MIN_LOSS:
50         MIN_LOSS = total_loss
51         best_model = model.state_dict()
52         # Save the best model
53         if best_model is not None:
54             print('Saving Best Model: ', MODEL_SAVED)
55             torch.save(best_model, MODEL_SAVED)
56
57 #####
58

```



Epoch 1/4, Loss: 0.4652398204101301

TRAINING DATA

- Performance:

	precision	recall	f1-score	support
0.0	0.88	0.90	0.89	7477
1.0	0.90	0.87	0.89	7523
accuracy			0.89	15000
macro avg	0.89	0.89	0.89	15000
weighted avg	0.89	0.89	0.89	15000

Saving checkpoint: results/checkpoints/epoch_1_20250619072045.pth

Saving Best Model: results/bestmodel.pth

Epoch 2/4, Loss: 0.18003429818366254

TRAINING DATA

- Performance:

	precision	recall	f1-score	support
0.0	0.92	0.98	0.95	7466
1.0	0.98	0.91	0.94	7534
accuracy			0.94	15000
macro avg	0.95	0.94	0.94	15000
weighted avg	0.95	0.94	0.94	15000

Saving checkpoint: results/checkpoints/epoch_2_20250619072407.pth

Saving Best Model: results/bestmodel.pth

Epoch 3/4, Loss: 0.12192414922000312

TRAINING DATA

- Performance:

	precision	recall	f1-score	support
0.0	0.97	0.95	0.96	7550
1.0	0.95	0.97	0.96	7450
accuracy			0.96	15000
macro avg	0.96	0.96	0.96	15000

```
weighted avg      0.96      0.96      0.96      15000
```

```
Saving checkpoint: results/checkpoints/epoch_3_20250619072732.pth
```

```
Saving Best Model: results/bestmodel.pth
```

```
-----
Epoch 4/4, Loss: 0.05339157602428071
```

```
TRAINING DATA
```

```
- Performance:
```

	precision	recall	f1-score	support
0.0	0.99	1.00	0.99	7483
1.0	1.00	0.99	0.99	7517
accuracy			0.99	15000
macro avg	0.99	0.99	0.99	15000
weighted avg	0.99	0.99	0.99	15000

```
1 model.load_state_dict(torch.load(MODEL_SAVED, weights_only
```

```
⇒ <All keys matched successfully>
```

✓ RESULTS ON ORIGINAL TRAINING AND TEST DATA

```
1 data_loader = create_training_loader() # without upsampling
```

```
2
```

```
3 # getting Best model performance on training, validation an
```

```
4 print('TRAINING DATA')
```

```
5 print(f'- Performance: \n{calculate_classification_accuracy
```

```
⇒ ::: TRAINING DATA DETAILS :::
```

```
- Number of Samples: 15000
```

```
- LABEL DISTRIBUTION:
```

```
  LABEL
```

```
1      10000
```

```
0       5000
```

```
Name: count, dtype: int64
```

```
TRAINING DATA
```

```
- Performance:
```

	precision	recall	f1-score	support
0.0	0.98	1.00	0.99	5000
1.0	1.00	0.99	0.99	10000
accuracy			0.99	15000
macro avg	0.99	0.99	0.99	15000
weighted avg	0.99	0.99	0.99	15000

```
1 testset_loader = create_training_loader(data_csv = TESTING_
```

```
2
```

```
3 print('TESTING DATA')
```

```
4 print(f'- Performance: \n{calculate_classification_accuracy
5
```



::: TRAINING DATA DETAILS :::

- Number of Samples: 3000

- LABEL DISTRIBUTION:

 LABEL

1 2000

0 1000

Name: count, dtype: int64

TESTING DATA

- Performance:

	precision	recall	f1-score	support
0.0	0.90	0.88	0.89	1000
1.0	0.94	0.95	0.95	2000
accuracy			0.93	3000
macro avg	0.92	0.92	0.92	3000
weighted avg	0.93	0.93	0.93	3000